

R cookbook for the casual dabbler (2nd edition)

Charles Coverdale

2025-02-12

Introduction

G'day and welcome to *R cookbook for the casual dabbler (2nd edition)*.

RCCD 1st edition was originally published in 2020 as a side project during the COVID-19 pandemic in Melbourne.

As I wrote in the midst of lockdown:

I use R a lot for work and for side projects. Over the years I've collated a bunch of useful scripts, from macroeconomic analysis to quick hacks for making map legends format properly.

Historically my code has been stored in random Rpubs documents, medium articles, and a bunch of .Rmd files on my harddrive. Occasionally I feel like doing things properly - and upload code to a repository on github.

It doesn't take a genius to realize this isn't a very sustainable solution - and it also isn't very useful for sharing code with others. It turns out 2-years of lockdown in Melbourne was enough incentive to sit down and collate my best and most useful code into a single place. In the spirit of open source, a book seemed like the most logical format. The following is a very rough book written in **markdown** - R's very own publishing language.

RCCD1e had surprisingly good (and long-lasting reviews). Alas, five-years on, a lot has changed - in both the R community, and the world more broadly. As such, RCCD2e includes significant revisions. The most major of these is a restructure to make the chapters flow more logically. The Australian specific chapters have also been grouped together (e.g. election data).

Usage

In each chapter I've written up the background, methodology and code for a separate piece of analysis.

Most of this code will not be extraordinary to the seasoned R aficionado. However I find that in classic Pareto style ~20% of my code contributes to the **vast** majority of my work output. Having this 20% on hand will hopefully be useful to both myself and others.

Additional resources

The R community is continually writing new packages and tools. Many of these are covered extensively in various free books available on the bookdown.org website.

The rise of LLM's over the past 2-years has also made it significantly easier to find, refine, and expand on R code. I recommend using them as a first point of call.

Limitations

If you find a bug (along with spelling errors etc) please email me at charlesf-coverdale@gmail.com with the subject line 'RCCD2e'.

About the author

Charles Coverdale is an economist based in Melbourne, Australia. He is passionate about economics, climate science, and building talented teams. You can get in touch with Charles on twitter to hear more about his current projects.

Making maps beautiful

Why use a map

Maps are a great way to communicate data.

They're easily understandable, flexible, and more intuitive than a chart. There's been numerous studies showing that the average reader often struggles to interpret the units on a y-axis, let alone understand trends in scatter or line graphs.

Making maps in R takes some initial investment. However once you have some code you know and understand, spinning up new pieces of analysis can happen in minutes, rather than hours or days.

The aim of this chapter is to get you from 'I can make a map in R' to something more like 'I can conduct spatial analysis and produce a visual which is ready for publication'.

Getting started

First up, we need to load a bunch of mapping packages. The tidyverse package is a classic for just about everything data manipulation, while, the **sf** and **ggspatial** packages are essential for making maps.

```
# Load required packages
library(tidyverse) # Includes ggplot2, dplyr, tidyr, readxl, purrr
library(ggmap)
library(sf)
library(ggspatial)
library(rlang)
library(broom)
library(Census2016)
library(strayr)
library(absmapsdata)
library(officer)
```

Will Mackey's `absmapsdata` package contains all the ABS' ASGS shapefiles. The data is now also callable through the (thoroughly updated and expanded) `strayr` package. For example:

```
strayr::read_absmmap("sa12016")
```

To get a basic demographic map up and running, we will splice together the shapefile (in this case the SA2 map of Australia) and some data from the 2016 Australian Census.

Hugh Parsonage put together a fantastic package called `Census2016` which makes downloading this data in a clean format easy.

```
#Get the shapefile form the absmapsdata package (predefined in the list above)

#Get the 2016 census dataset
census2016_wide <- Census2016_wide_by_SA2_year

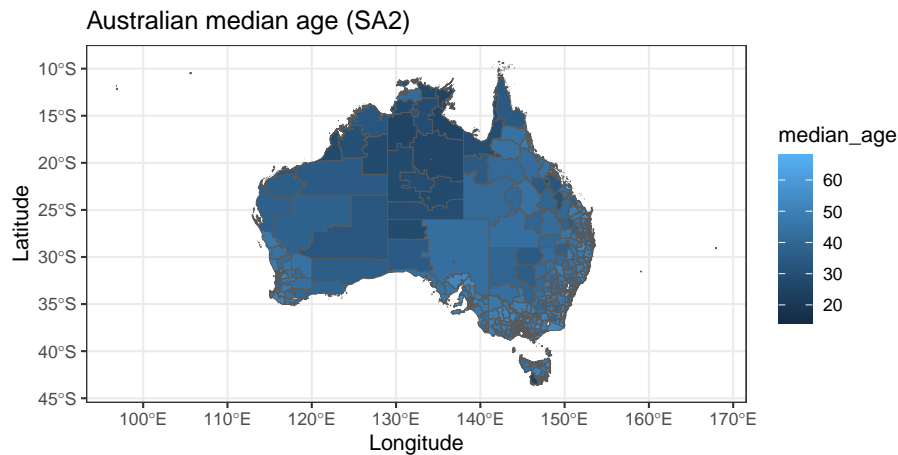
#Select the key demographic columns from the census data (i.e. the first 8 variables)
census_short <- census2016_wide[,1:8]

#Filter for a single year
census_short_2016 <- census_short %>%
  filter(year==2016)

#Use the inner_join function to get the shapefile and census wide data into a single d.
SA2_shp_census_2016 <- inner_join(strayr::read_absmmap("sa22016"),census_short_2016,
                                by = c("sa2_name_2016" = "sa2_name"))

#Plot a map that uses census data
map1 <- ggplot() +
  geom_sf(data = SA2_shp_census_2016, aes(fill = median_age)) +
  ggtitle("Australian median age (SA2)") +
  xlab("Longitude") +
  ylab("Latitude") +
  theme_bw() +
  theme(legend.position = "right")

map1
```



There we go! A map. This looks ‘okay’... but it can be much better.

From okay to good

Heat maps don’t really show too much interesting data on such a large scale, so let’s filter down to Greater Melbourne.

Seeing we have a bunch of census data in our dataframe, we can also do some basic analysis (e.g. population density).

```
#As a bit of an added extra, we can create a new population density column
SA2_shp_census_2016 <- SA2_shp_census_2016 %>%
  mutate(pop_density=persons/areasqkm_2016)

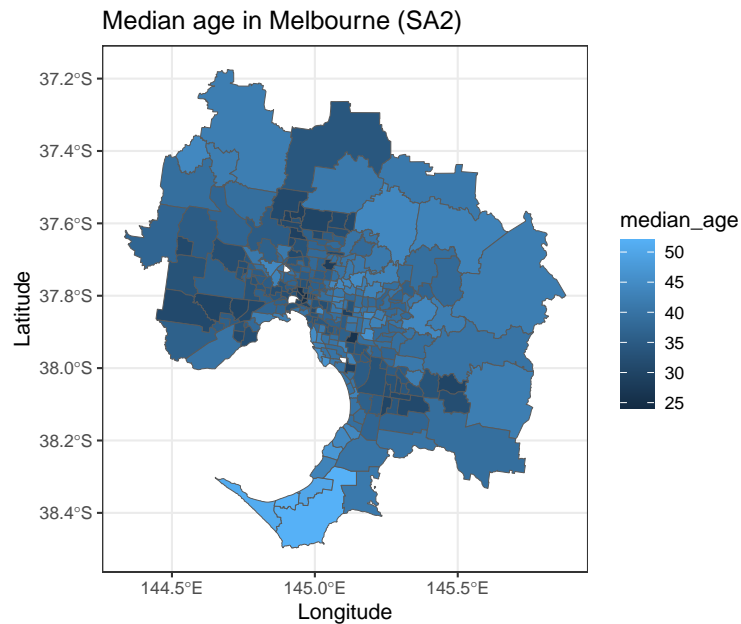
#Filter for Greater Melbourne
MEL_SA2_shp_census_2016 <- SA2_shp_census_2016 %>%
  filter(gcc_name_2016=="Greater Melbourne")

#Plot the new map just for Greater Melbourne
map2 <- ggplot() +
  geom_sf(data = MEL_SA2_shp_census_2016, aes(fill = median_age, border=NA)) +
  ggtitle("Median age in Melbourne (SA2)") +
  xlab("Longitude") +
  ylab("Latitude") +
```

```

theme_bw() +
theme(legend.position = "right")
map2

```



Much better. We can start to see some trends in this map. It looks like younger people tend to live closer to the city center. This seems logical.

From good to great

The map above is a good start! However, how do we turn this from something ‘good’, into something that is 100% ready to share?

We see our ‘ink to chart ratio’ (i.e. the amount of non-data stuff that is on the page) is still pretty high. Is the latitude of Melbourne useful for this analysis...? Not really. Let’s get rid of it and the axis labels. A few lines of code adjusting the axis, titles, and theme of the plot will go a long way. Because my geography Professor drilled it into me, I will also add a low-key scale bar.

```

map3 <- ggplot() +
  geom_sf(data = MEL_SA2_shp_census_2016, aes(fill = median_age)) +
  labs(title="Melbourne's youth tend to live closer to the city centre",
        subtitle = "Analysis from the 2016 census",
        caption = "Data: Australian Bureau of Statistics 2016",

```



```

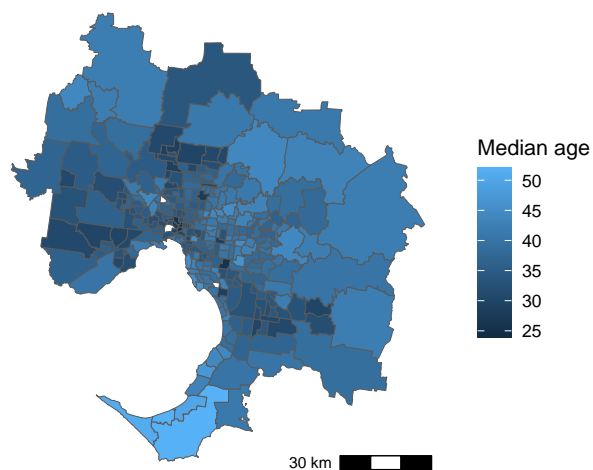
x="",
y="",
fill="Median age") +
ggspatial::annotation_scale(location="br")+
theme_minimal() +
  theme(axis.ticks.x = element_blank(),axis.text.x = element_blank())+
  theme(axis.ticks.y = element_blank(),axis.text.y = element_blank())+
  theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank())+
  theme(legend.position = "right")+
  theme(plot.title=element_text(face="bold",size=12))+
  theme(plot.subtitle=element_text(size=11))+
  theme(plot.caption=element_text(size=8))

```

map3

Melbourne's youth tend to live closer to the city centre

Analysis from the 2016 census



From great to fantastic

The above is perfectly reasonable and looks professionally designed. However, this is where we can get really special.

Let's add a custom colour scheme, drop the boundary edges for the SA2's, and add in a dot and label for Melbourne CBD.

```

#Add in a point for the Melbourne CBD
MEL_location <- data.frame(town_name = c("Melbourne"),
                           x = c(144.9631),
                           y = c(-37.8136))

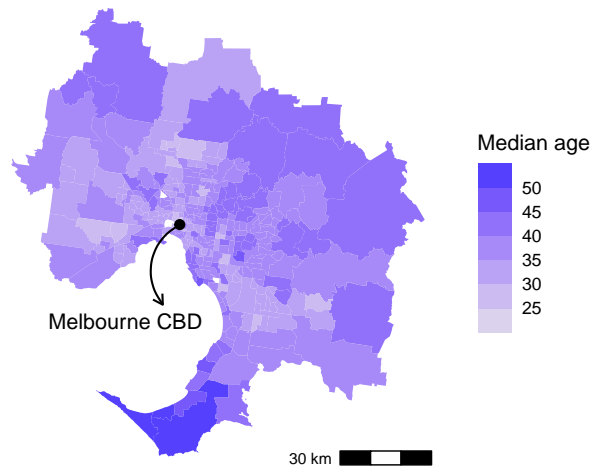
map4 <- ggplot() +
  geom_sf(data = MEL_SA2_shp_census_2016, aes(fill = median_age), color=NA) +
  geom_point(data=MEL_location, aes(x=x, y=y), size=2, color="black") +
  labs(title="Melbourne's youth tend to live closer to the city centre",
       subtitle = "Analysis from the 2016 census",
       caption = "Data: Australian Bureau of Statistics 2016",
       x="",
       y="",
       fill="Median age") +
  scale_fill_steps(low="#E2E0EB", high="#3C33FE") +
  annotate(geom='curve',
          x=144.9631,
          y=-37.8136,
          xend=144.9,
          yend=-38.05,
          curvature=0.5,
          arrow=arrow(length=unit(2, "mm")))+
  annotate(geom='text', x=144.76, y=-38.1, label="Melbourne CBD")+
  ggspatial::annotation_scale(location="br")+
  theme_minimal() +
  theme(axis.ticks.x = element_blank(), axis.text.x = element_blank())+
  theme(axis.ticks.y = element_blank(), axis.text.y = element_blank())+
  theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank())+
  theme(legend.position = "right")+
  theme(plot.title=element_text(face="bold", size=12))+
  theme(plot.subtitle=element_text(size=11))+
  theme(plot.caption=element_text(size=8))

map4

```

Melbourne's youth tend to live closer to the city centre

Analysis from the 2016 census



Data: Australian Bureau of Statistics 2016

Now we're talking. A 'client-ready' looking map that can be added to a report, presentation, or with a few tweaks - a digital dashboard.

Make sure to export the map as a high quality PNG using the `ggplot2::ggsave()` function.

Charts

Getting started

There's exceptional resources online for using the `ggplot2` package to create production ready charts.

The R Graph Gallery is a great place to start, as is the visual storytelling blogs of The Economist and the BBC.

This chapter contains the code for some of my most used charts and visualization techniques.

```
# Load in packages
library(ggribes)
library(ggplot2)
library(ggrepel)
library(viridis)
library(readxl)
library(dplyr)
library(stringr)
library(reshape)
library(tidyr)
library(lubridate)
library(gapminder)
#library(grattantheme)
library(patchwork)
library(scales)
library(purrr)
```

Make the data tidy

Before making a chart ensure the data is “tidy” - meaning there is a new row for every changed variable. It also doesn't hurt to remove NA's for consistency (particularly in time series).

```

#Read in data
url <- "https://raw.githubusercontent.com/charlescoverdale/ggbridges/master/2019_MEL_max"

#Read in with read.xlsx
MEL_temp_daily <- openxlsx::read.xlsx(url)

#Remove last 2 characters to just be left with the day number
MEL_temp_daily$Day=substr(MEL_temp_daily$Day,1,nchar(MEL_temp_daily$Day)-2)

#Make a wide format long using the gather function
MEL_temp_daily <- MEL_temp_daily %>%
  gather(Month,Temp,Jan:Dec)

MEL_temp_daily$Month<-factor(MEL_temp_daily$Month,levels=c("Jan", "Feb", "Mar", "Apr",

#Add in a year
MEL_temp_daily["Year"]=2019

#Reorder
MEL_temp_daily <- MEL_temp_daily[,c(1,2,4,3)]

#Make a single data field using lubridate
MEL_temp_daily <- MEL_temp_daily %>% mutate(Date = make_date(Year, Month, Day))

#Drop the original date columns
MEL_temp_daily <- MEL_temp_daily %>% dplyr::select(Date, Temp) %>% drop_na()

#Add on a 7-day rolling average
MEL_temp_daily <- MEL_temp_daily %>% dplyr::mutate(Seven_day_rolling =
  zoo::rollmean(Temp, k = 7, fill = NA,
  Mean = mean(Temp))

#Drop NA's
#MEL_temp_daily <- MEL_temp_daily %>% drop_na()

```

Line plot

```

plot_MEL_temp <- ggplot(MEL_temp_daily)+
  geom_line(aes(x = Date, y = Temp), col = "blue")+
  geom_line(aes(x = Date, y = Mean), col = "orange")+
  labs(title="Hot in the summer and cool in the winter",
  subtitle = "Analysing temperature in Melbourne",
  caption = "Data: Bureau of Meterology 2019",

```

```

x="",
y="") +

scale_x_date(date_breaks = "1 month",
             date_labels = "%b",
             limits = as.Date(c('2019-01-01', '2019-12-14')))+

scale_y_continuous(label = unit_format(unit="\u00b0C", sep="")) +

theme_minimal() +
theme(legend.position="bottom")+

theme(plot.title=element_text(face="bold",size=12))+
theme(plot.subtitle=element_text(size=11))+
theme(plot.caption=element_text(size=8))+

theme(axis.text=element_text(size=8))+
theme(panel.grid.minor = element_blank())+
theme(panel.grid.major.x = element_blank()) +

theme(axis.title.y =
      element_text(margin = margin(t = 0, r = 0, b = 0, l = 0)))+

theme(axis.text.y = element_text(vjust = -0.5,
                                  margin = margin(l = 20, r = -20)))+

theme(axis.line.x = element_line(colour = "black",size=0.4))+
theme(axis.ticks.x = element_line(colour = "black",size=0.4))+


annotate(geom='curve',
         x=as.Date('2019-08-01'), y=23,
         xend=as.Date('2019-08-01'),yend=17,
         curvature=-0.5,arrow=arrow(length=unit(2,"mm")))+
annotate(geom='text',x=as.Date('2019-07-15'),y=25,
         label="Below 20°C all winter")

```

plot_MEL_temp

Hot in the summer and cool in the winter

Analysing temperature in Melbourne



Data: Bureau of Meteorology 2019

Scatter and trend plot

```
MEL_temp_Jan <- MEL_temp_daily %>% filter(MEL_temp_daily$Date<as.Date("2019-01-31"))

ggplot(MEL_temp_Jan)+
  geom_point(aes(x = Date, y = Temp), col = "purple",alpha=0.4)+
  geom_smooth(aes(x = Date, y = Temp), col = "purple",fill="purple", alpha=0.1,method = "loess",
    #scale_colour_manual(values = c("purple"),labels="Trend") +
    #scale_fill_manual(values = c("purple"),labels="Confidence interval")+
  labs(title="January is a hot one",
    subtitle = "Analysing temperature in Melbourne",
    caption = "Data: Bureau of Meterology 2019",
    x="",
    y="Temperature °C") +

  scale_x_date(date_breaks = "1 week",
    date_labels = "%d-%b",
    limits = as.Date(c('2019-01-01','2019-01-31')),
    expand = c(0,0)) +

  theme_minimal() +
```



```

theme(plot.title=element_text(face="bold",size=12))+
theme(plot.subtitle=element_text(size=11))+
theme(plot.caption=element_text(size=8))+

theme(axis.text=element_text(size=8))+
theme(axis.title.y = element_text(size=9,margin=margin(t = 0, r = 10, b = 0, l = 0)))+
theme(panel.grid.minor = element_blank()+
theme(panel.grid.major.x = element_blank()) +

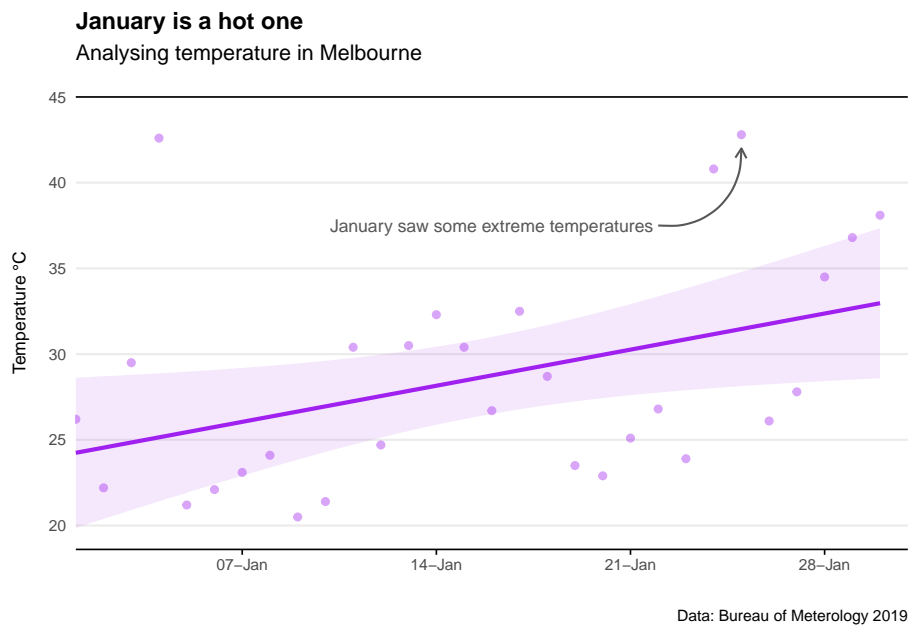
theme(axis.line.x = element_line(colour ="black",size=0.4))+
theme(axis.ticks.x = element_line(colour ="black",size=0.4))+

geom_hline(yintercept=45,colour ="black",size=0.4) +

annotate(geom='curve',
x=as.Date('2019-01-22'), y=37.5,
xend=as.Date('2019-01-25'),yend=42,
curvature=0.5,
col="#575757",
arrow=arrow(length=unit(2,"mm")))) +

annotate(geom='text',x=as.Date('2019-01-16'),y=37.5,
label="January saw some extreme temperatures",size=3.2,col="#575757")

```



Shading areas on plots

Adding shading behind a plot area is simple using `geom_rect`.

Adding shading under a particular model line? A little trickier. See both example below.


```
#Example 1
#Set a custom limit for the y-axis threshold
threshold <- 20

ggplot() +
  geom_point(data = mtcars,
            aes(x = hp, y = mpg)) +

  #Add an intercept line
  geom_hline(yintercept = threshold) +

  # Shade area under y_lim
  geom_rect(aes(xmin = -Inf, xmax = Inf, ymin = -Inf, ymax = threshold),
            alpha = 1/5,
            fill = "blue") +

  # Shade area above y_lim
  geom_rect(aes(xmin = -Inf, xmax = Inf, ymin = threshold, ymax = Inf),
            alpha = 1/5,
            fill = "red")
```



```
#Example 2
#Now to add colour either way of a model line

#Define the model
model <- lm(mpg ~ log(hp), data = mtcars)

# Get the predictions for plotting. Here, df_line, is a data frame with new
# coordinates that will be used for plotting the trend line and further for
# building the polygons for shading.
min_x <- min(mtcars$hp)
max_x <- max(mtcars$hp)

df_line <- data.frame(hp = seq(from = min_x, to = max_x, by = 1))
df_line$mpg <- predict(model, newdata = df_line)

p <- ggplot() +
  geom_point(data = mtcars,
             aes(x = hp, y = mpg), color = "grey", alpha = 0.5) +

  geom_line(data = df_line,
            aes(x = hp, y = mpg), col = "black")

#Define two polygons (one above, and one below the line)
df_poly_under <- df_line %>%
  tibble::add_row(hp = c(max_x, min_x),
```

```

      mpg = c(-Inf, -Inf))

df_poly_above <- df_line %>%
  tibble::add_row(hp = c(max_x, min_x),
                 mpg = c(Inf, Inf))

#Plot the data, line, and the shades above and below the line
p +

  geom_polygon(data = df_poly_under,
              aes(x = hp, y = mpg),
              fill = "blue",
              alpha = 1/5) +

  geom_polygon(data = df_poly_above,
              aes(x = hp, y = mpg),
              fill = "red",
              alpha = 1/5)+

  scale_x_discrete(expand = c(0,0))+

  theme_minimal() +

  labs(title="Look at that snazzy red/blue shaded area",
       subtitle = "Subtitle goes here",
       caption = "Data: Made up from scratch",
       x="",
       y="") +

  theme(panel.grid.minor = element_blank())+
  theme(panel.grid.major.x = element_blank()) +


  theme(axis.title.y=element_blank(),
        axis.text.y=element_blank(),
        axis.ticks.y=element_blank())+

  theme(axis.title.x=element_blank(),
        axis.text.x=element_blank(),
        axis.ticks.x=element_blank())

```

Look at that snazzy red/blue shaded area

Subtitle goes here



Data: Made up from scratch

Bar chart (numercial)

```
Year = c("2018", "2019", "2020", "2021")

Value = (c(1000000, 3000000, 2000000, 5000000))

bar_data_single <- (cbind(Year, Value))

bar_data_single <- as.data.frame(bar_data_single)

bar_data_single$Value = as.integer(bar_data_single$Value)

ggplot(bar_data_single, aes(x = Year, y = Value, label=Value)) +
  geom_bar(stat='identity', fill="blue", width=0.8) +
  geom_text(size = 5,
            col="white", fontface="bold",
            position = position_stack(vjust = 0.5),
            label=scales::dollar(Value, scale=1/1e6, suffix="m")) +

  labs(title="Bar chart example",
        subtitle = "Subtitle goes here",
```

```

caption = "Data: Made up from scratch",
x="",
y="") +

theme_minimal() +

theme(plot.title=element_text(face="bold",size=12))+
theme(plot.subtitle=element_text(size=11))+
theme(plot.caption=element_text(size=12))+


theme(axis.text=element_text(size=12))+
theme(panel.grid.minor = element_blank())+
theme(panel.grid.major.x = element_blank())+
theme(panel.grid.major.y = element_blank()) +

theme(axis.title.y=element_blank(),
      axis.text.y=element_blank(),
      axis.ticks.y=element_blank())

```

Bar chart example

Subtitle goes here



Data: Made up from scratch

```

# ggsave(plot=last_plot(),
#       width=10,
#       height=10,
#       units="cm",
#       dpi = 600,

```

```
# filename = "/Users/charlescoverdale/Desktop/test.png")
```

Stacked bar chart

```
Year = c("2019", "2019", "2019", "2019", "2020", "2020", "2020", "2020")

Quarter = c("Q1", "Q2", "Q3", "Q4", "Q1", "Q2", "Q3", "Q4")

Value = (c(100, 300, 200, 500, 400, 700, 200, 300))

bar_data <- (cbind(Year, Quarter, Value))

bar_data <- as.data.frame(bar_data)

bar_data$Value = as.integer(bar_data$Value)

bar_data_totals <- bar_data %>%
  dplyr::group_by(Year) %>%
  dplyr::summarise(Total = sum(Value))

ggplot(bar_data, aes(x = Year, y = Value, fill = (Quarter), label=Value)) +
  geom_bar(position = position_stack(reverse=TRUE), stat='identity') +
  geom_text(size = 4,
            col="white",
            fontface="bold",
            position = position_stack(reverse=TRUE, vjust = 0.5),
            label=scales::dollar(Value)) +
  geom_text(aes(Year, Total,
                label=scales::dollar(Total),
                fill = NULL,
                vjust=-0.5),
            fontface="bold",
            size=4,
            data = bar_data_totals) +

  scale_fill_brewer(palette = "Blues") +

  labs(title="Bar chart example",
        subtitle = "Subtitle goes here",
        caption = "Data: Made up from scratch",
        x="",
```

```

y="Units") +

theme_minimal() +

theme(legend.position = "bottom")+
theme(legend.title = element_blank())+

theme(plot.title=element_text(face="bold",size=12))+
theme(plot.subtitle=element_text(size=10))+
theme(plot.caption=element_text(size=8))+

theme(axis.text=element_text(size=10))+
theme(panel.grid.minor = element_blank())+
theme(panel.grid.major.x = element_blank())+
theme(panel.grid.major.y = element_blank()) +


scale_y_continuous(expand=c(0,0),limits=c(0,1800))+

theme(axis.title.y=element_blank(),
      axis.text.y=element_blank(),
      axis.ticks.y=element_blank())

```

Bar chart example

Subtitle goes here



Data: Made up from scratch

Histogram

Aka. a bar chart for a continuous variable where the bars are touching. Useful to show distribution of time series or ordinal variables.

```
c("0-20", "20-40", "40-60", "60-80", "80+")
```

```
## [1] "0-20" "20-40" "40-60" "60-80" "80+"
```

```
#Create a data set
hist_data <- data.frame(X1=sample(0:100,100,rep=TRUE))

ggplot(hist_data)+

  geom_histogram(aes(x=X1),binwidth=5,fill="blue",alpha=0.5) +

  geom_vline(xintercept=c(50,75,95),yintercept=0,linetype="longdash",col="orange") +

  labs(title="Histogram example",
        subtitle = "Facet wraps are looking good",
        caption = "Data: Made up from scratch",
        x="",
        y="") +

  theme_minimal() +

  theme(panel.spacing.x = unit(10, "mm"))+

  theme(legend.position="none")+

  theme(plot.title=element_text(face="bold",size=12))+
  theme(plot.subtitle=element_text(size=11))+
  theme(plot.caption=element_text(size=8))+


  theme(axis.text=element_text(size=9))+
  theme(panel.grid.minor = element_blank())+
  theme(panel.grid.major.x = element_blank()) +

  theme(plot.subtitle = element_text(margin=margin(0,0,15,0))) +

  theme(plot.margin=unit(c(0.5,0.5,0.5,0.5),"cm"))
```

Histogram example

Facet wraps are looking good

**Ridge chart**

Handy when working with climate variables. Particularly useful at showing the difference in range of multiples series (e.g. temperature by month).

```
# Import data
url <- "https://raw.githubusercontent.com/charlescoverdale/ggribges/master/2019_MEL_max"

MEL_temp_daily <- openxlsx::read.xlsx(url)

# Remove last 2 characters to just be left with the day number
MEL_temp_daily$Day = substr(MEL_temp_daily$Day, 1, nchar(MEL_temp_daily$Day) - 2)


# Make a wide format long using the gather function
MEL_temp_daily <- MEL_temp_daily %>%
  gather(Month, Temp, Jan:Dec)

MEL_temp_daily$Month <- factor(MEL_temp_daily$Month, levels = c("Jan", "Feb", "Mar", "Apr",
# Plot
ggplot(MEL_temp_daily,
      aes(x = Temp, y = Month, fill = stat(x))) +
  geom_density_ridges_gradient(scale = 2,
```

```

size=0.3,
rel_min_height = 0.01,
gradient_lwd = 1.) +
scale_y_discrete(limits = unique(rev(MEL_temp_daily$Month)))+
scale_fill_viridis_c(name = "°C", option = "C") +
labs(title = 'Melbourne temperature profile',
      subtitle = 'Daily maximum temperature recorded in Melbourne in 2019',
      caption = "Data: Bureau of Meteorology 2020") +
xlab(" ") +
ylab(" ") +
theme_ridges(font_size = 13, grid = TRUE)

```



BBC style: Bar charts (categorical)

```

#devtools::install_github('bbc/bbplot')
library(gapminder)
library(bbplot)

# Prepare data
bar_df <- gapminder %>%
  filter(year == 2007 & continent == "Africa") %>%

```

```


  arrange(desc(lifeExp)) %>%
  head(5)

# Make plot
bars <- ggplot(bar_df, aes(x = reorder(country, lifeExp), y = lifeExp, fill = country)) +
  geom_bar(stat="identity", position="identity") +
  geom_hline(yintercept = 0, size = 1, colour="#333333") +
  scale_fill_manual(values = c("TRUE" = "#1380A1", "FALSE" = "#ddddd")) +
  labs(title = "Mauritius has the highest life expectancy",
       subtitle = "Top 5 African countries by life expectancy, 2007") +
  coord_flip() +
  theme_minimal() +
  theme(panel.grid.major.x = element_line(color="#cbcbcb"),
        panel.grid.major.y = element_blank(),
        legend.position = "none")

# Add labels
labelled_bars <- bars +
  geom_label(aes(label = round(lifeExp, 0)),
            hjust = 1, vjust = 0.5, colour = "white",
            fill = "black", label.size = 0.2,
            family = "Helvetica", size = 6)

labelled_bars

```



BBC style: Dumbbell charts


Dumbbell charts are handy instead of using clustered column charts with janky thinkcell labels and arrows to show the difference between the columns. Note it relies on having a 4 variable input (variable_name, value1, value2, and gap).

The `geom_dumbbell` function lives inside the `ggalt` package rather than the standard `ggplot2`.

```
library(ggalt)    # For geom_dumbbell
library(gapminder)
library(bbbplot)  # Uncomment if using bbc_style()

# Prepare data
dumbbell_df <- gapminder %>%
  filter(year %in% c(1967, 2007)) %>%
  select(country, year, lifeExp) %>%
  pivot_wider(names_from = year, values_from = lifeExp) %>%
  mutate(gap = `2007` - `1967`) %>%
  arrange(desc(gap)) %>%
  head(10)

# Make plot
ggplot(dumbbell_df, aes(x = `1967`, xend = `2007`, y = reorder(country, gap), group = country)) +
  geom_dumbbell(colour = "#ddddd",
                size = 3,
                colour_x = "#FAAB18",
                colour_xend = "#1380A1") +
  labs(title = "We're Living Longer",
        subtitle = "Top 10 biggest increases in life expectancy (1967-2007)") +
  theme_minimal() +
  theme(panel.grid.major.y = element_blank(),
        panel.grid.minor = element_blank())
```



Facet wraps

Handy rather than showing multiple lines on the same chart.

Top tips: `facet_wrap()` dataframes need to be in long form in order to be manipulated easily.

It also helps to add on separate columns for the start and end values (if you want to add data point labels).

```
#Create a data set
Year = c("2018", "2019", "2020", "2021")

QLD = (c(500,300, 500, 600))

NSW = (c(200,400, 500, 700))

VIC = (c(300,400, 500, 600))

#Combine the columns into a single dataframe
facet_data <- (cbind(Year, QLD,NSW,VIC))
facet_data <- as.data.frame(facet_data)

#Change formats to integers
```

```

facet_data$QLD = as.integer(facet_data$QLD)
facet_data$NSW = as.integer(facet_data$NSW)
facet_data$VIC = as.integer(facet_data$VIC)

#Make the wide data long
facet_data_long <- pivot_longer(facet_data,!Year, names_to="State", values_to="Value")

facet_data_long <- facet_data_long %>%
  dplyr::mutate(start_label =
    if_else(Year == min(Year),
            as.integer(Value), NA_integer_))

facet_data_long <- facet_data_long %>%
  dplyr::mutate(end_label =
    if_else(Year == max(Year),
            as.integer(Value), NA_integer_))

#Make the base line chart
base_chart <- ggplot() +

  geom_line(data=facet_data_long,
    aes(x = Year,
        y = Value,
        group = State,
        colour = State)) +

  geom_point(data=facet_data_long,
    aes(x = Year,
        y = Value,
        group = State,
        colour = State)) +

  ggrepel::geom_text_repel(data=facet_data_long,
    aes(x = Year,
        y = Value,
        label = end_label),
    color = "black",
    nudge_y = -10,size=3) +

  ggrepel::geom_text_repel(data=facet_data_long,
    aes(x = Year,
        y = Value,
        label = start_label),

```

```

        color = "black",
        nudge_y = 10,size=3)

base_chart +

    scale_x_discrete(
        breaks = seq(2018, 2021, 1),
        labels = c("2018", "19", "20", "21"))+

    facet_wrap(State ~ .) +
        #To control the grid arrangement, we can add in customer dimensions
        #ncol = 2, nrow=2) +

    labs(title="State by state comparison",
          subtitle = "Facet wraps are looking good",
          caption = "Data: Made up from scratch",
          x="",
          y="") +

    theme_minimal() +

    theme(strip.text.x = element_text(size = 9, face = "bold")) +

    theme(panel.spacing.x = unit(10, "mm"))+

    theme(legend.position="none")+


    theme(plot.title=element_text(face="bold",size=12))+
    theme(plot.subtitle=element_text(size=11))+
    theme(plot.caption=element_text(size=8))+

    theme(axis.text=element_text(size=9))+
    theme(panel.grid.minor = element_blank())+
    theme(panel.grid.major.x = element_blank()) +

    theme(plot.subtitle = element_text(margin=margin(0,0,15,0))) +

    theme(plot.margin=unit(c(0.5,0.5,0.5,0.5),"cm"))

```

Pie chart

These should be used sparingly... but they are handy for showing proportions when the proportion of the whole is paramount (e.g. 45%) - rather than the proportion in relation to another data point (e.g. 16% one year vs 18% the next).

```
# Create Data
pie_data <- data.frame(
  group=LETTERS[1:5],
  value=c(13,7,9,21,2))

# Compute the position of labels
pie_data <- pie_data %>%
  arrange(desc(group)) %>%
  mutate(proportion = value / sum(pie_data$value) *100) %>%
  mutate(ypos = cumsum(proportion)- 0.5*proportion )

# Basic piechart
ggplot(pie_data, aes(x="", y=proportion, fill=group)) +
  geom_bar(stat="identity")+
  coord_polar("y", start=0) +
  theme_void()+
```

```

geom_text(aes(y = ypos,
              label = paste(round(proportion,digits=0),"%", sep = ""),x=1.25),
          color = "white",
          size=4) +

scale_fill_brewer(palette="Set1")+

labs(title="Use pie charts sparingly",
      subtitle = "Subtitle goes here",
      caption = "",
      x="",
      y="") +


theme(legend.position = "bottom")+
theme(legend.title = element_blank())+

theme(plot.title=element_text(face="bold",size=12))+
theme(plot.subtitle=element_text(size=11))+
theme(plot.caption=element_text(size=8))+

theme(plot.subtitle = element_text(margin=margin(0,0,5,0))) +
theme(plot.margin=unit(c(0.5,0.5,0.5,0.5),"cm"))

```

Use pie charts sparingly
 Subtitle goes here




```
# ggsave(plot=last_plot(),  
#       width=10,  
#       height=10,  
#       units="cm",  
#       dpi = 600,  
#       filename = "/Users/charlescoverdale/Desktop/pietest.png")
```


Patchwork

Patchwork is a nifty package for arranging plots and other graphic elements (text, tables etc) in different grid arrangements. The basic syntax is to use `plot1 | plot2` for side by side charts, and `plot1 / plot2` for top and bottom charts. You can also combine these two functions for a grid of different size columns (e.g. `plot3 / (plot1 | plot2)`)

```
#Make some simply plots using the mtcars package  
p1 <- ggplot(mtcars) +  
  geom_point(aes(mpg, disp)) +  
  ggtitle('Plot 1')  
  
p2 <- ggplot(mtcars) +  
  geom_boxplot(aes(gear, disp, group = gear)) +  
  ggtitle('Plot 2')  
  
#Example of side by side charts  
library(patchwork)  
p1 + p2
```



```
#Add in a table next to the plot
p1 + gridExtra::tableGrob(mtcars[1:10, c('mpg', 'disp')])
```



Saving to powerpoint

There's a bunch of ways to save ggplot graphics - but the way I find most useful is by exporting to pptx in a common 'charts' directory.

If you want to save as a png you can use the normal ggsave function - however it will not be editable (e.g. able to click and drag to rescale for a presentation).

Therefore instead we can use the grattantheme package to easily save to an editable pptx graphic.

Note: The code below has been commented out so that it will upload to book-down.org without an error.

```
#The classic save function to png

#   ggsave(plot = ggplot2::last_plot(),
#   width = 8,
#   height = 12,
#   dpi = 600,
#   filename = "/Users/charlescoverdale/Desktop/test.png")

#Using the grattantheme package to easily save to powerpoint

#   grattan_save_pptx(p = ggplot2::last_plot(),
#   "/Users/charlescoverdale/Desktop/test.pptx",
#   type = "wholecolumn")
```

Automating chart creation

Let's say we have a dataframe of multiple variables. We want to produce simple charts of the same style for each variable (including formatting and titles etc). Sure we can change the `aes(x=)` and `aes(y=)` variables in ggplot2 manually for each column - but this is time intensive especially for large data frames. Instead, we can write a function that will loop through the whole dataframe and produce the same format of chart.

```
#Create a data set
Year = c("2018", "2019", "2020", "2021")

Variable1 = (c(500,300, 200, 400))

Variable2 = (c(200,400, 200, 700))

Variable3 = (c(300,500, 800, 1000))
```

```

#Combine the columns into a single dataframe
bar_data_multiple <- (cbind(Year, Variable1, Variable2, Variable3))
bar_data_multiple <- as.data.frame(bar_data_multiple)

#Change formats to integers
bar_data_multiple$Variable1 = as.integer(bar_data_multiple$Variable1)
bar_data_multiple$Variable2 = as.integer(bar_data_multiple$Variable2)
bar_data_multiple$Variable3 = as.integer(bar_data_multiple$Variable3)

#Define a function
loop <- function(chart_variable) {

  ggplot(bar_data_multiple, aes(x = Year,
                                y = .data[[chart_variable]],
                                label = .data[[chart_variable]])) +

    geom_bar(stat='identity', fill="blue") +

    geom_text(aes(
      label = scales::dollar(.data[[chart_variable]]),
      size = 5,
      col="white",
      fontface="bold",
      position = position_stack(vjust = 0.5)) +

    labs(title=paste("Company X: ",
                     chart_variable,
                     " (", head(Year, n=1),
                     " - ",
                     tail(Year, n=1),
                     ")",
                     sep=""),
          subtitle = "Subtitle goes here",
          caption = "Data: Made up from scratch",
          x="",
          y="") +

    theme_minimal() +

    theme(plot.title=element_text(face="bold", size=12)) +
    theme(plot.subtitle=element_text(size=11)) +
    theme(plot.caption=element_text(size=12)) +

    theme(axis.text=element_text(size=12)) +

```


```
theme(panel.grid.minor = element_blank())+
theme(panel.grid.major.x = element_blank())+
theme(panel.grid.major.y = element_blank()) +

theme(axis.title.y=element_blank(),
      axis.text.y=element_blank(),
      axis.ticks.y=element_blank())
}

plots <- purrr::map(colnames(bar_data_multiple)[colnames(bar_data_multiple) != "Year"], loop)
plots
```

Company X: Variable1 (2018 – 2021)


Subtitle goes here



Data: Made up from scratch

Company X: Variable2 (2018 – 2021)


Subtitle goes here



Data: Made up from scratch

Company X: Variable3 (2018 – 2021)

Subtitle goes here



Data: Made up from scratch

```
#cowplot::plot_grid(plotlist = plots)
```


Basic modelling

Creating a model is an essential part of forecasting and data analysis.

I've put together a quick guide on my process for modelling data and checking model fit.

The source data I use in this example is Melbourne's weather record over a 12 month period. Daily temperature is based on macroscale weather and climate systems, however many observable measurements are correlated (i.e. hot days tend to have lots of sunshine). This makes using weather data great for model building.

Source, format, and plot data

Before we get started, it is useful to have some packages up and running.

```
#Useful packages for regression
library(readr)
library(readxl)
library(ggplot2)
library(tidyverse)
library(lubridate)
library(modelr)
library(cowplot)
```

I've put together a csv file of weather observations in Melbourne in 2019. We begin our model by downloading the data from Github.

```
#Input data
link <- "data/MEL_weather_2019.csv"

# We'll read this data in as a dataframe
# The 'check.names' function set to false means the funny units that the BOM use for column names
```

```
MEL_weather_2019 <- read.csv(link, check.names = F)

head(MEL_weather_2019)
```

This data is relatively clean. One handy change to make is to make the date into a dynamic format (to easily switch between months, years, etc).

```
#Add a proper date column
MEL_weather_2019 <- MEL_weather_2019 %>%
  mutate(Date = make_date(Year, Month, Day))
```

We also notice that some of the column names have symbols in them. This can be tricky to work with, so let's rename some columns into something more manageable.


```
#Rename key df variables
names(MEL_weather_2019)[4] <- "Solar_exposure"
names(MEL_weather_2019)[5] <- "Rainfall"
names(MEL_weather_2019)[6] <- "Max_temp"

head(MEL_weather_2019)
```

We're aiming to investigate if other weather variables can predict maximum temperatures. Solar exposure seems like a plausible place to start. We start by plotting the two variables to if there is a trend.

```
#Plot the data
MEL_temp_investigate <- ggplot(MEL_weather_2019)+
  geom_point(aes(y=Max_temp, x=Solar_exposure),col="grey")+
  labs(title = "Does solar exposure drive temperature in Melbourne?",
       caption = "Data: Bureau of Meteorology 2020") +
  xlab("Solar exposure")+
  ylab("Maximum temperature °C")+
  scale_x_continuous(expand=c(0,0))+
  theme_bw()+
  theme(axis.text=element_text(size=10))+
  theme(panel.grid.minor = element_blank())

MEL_temp_investigate
```




Eyeballing the chart above, there seems to be a correlation between the two data sets. We'll do one more quick plot to analyse the data. What is the distribution of temperature?

```
ggplot(MEL_weather_2019, aes(x=Max_temp)) +
  geom_histogram(aes(y=..density..), colour="black", fill="lightblue")+
  geom_density(alpha=.5, fill="grey", colour="darkblue")+

  scale_x_continuous(breaks=c(5,10,15,20,25,30,35,40,45),
                     expand=c(0,0))+
  xlab("Temperature")+
  ylab("Density")+

  theme_bw()+
  theme(axis.text=element_text(size=12))+
  theme(panel.grid.minor = element_blank())
```



We can see here the data is right skewed (i.e. the mean will be greater than the median). We'll need to keep this in mind. Let's start building a model.

Build a linear model

We start by looking whether a simple linear regression of solar exposure seems to be correlated with temperature. In R, we can use the linear model (`lm`) function.

```
#Create a straight line estimate to fit the data
temp_model <- lm(Max_temp~Solar_exposure, data=MEL_weather_2019)
```

Analyse the model fit

Let's see how well solar exposure explains changes in temperature

```
#Call a summary of the model
summary(temp_model)
```

The adjusted R squared value (one measure of model fit) is 0.3596. Furthermore the coefficient of our solar_exposure variable is statistically significant.

Compare the predicted values with the actual values

We can use this `lm` function to predict values of temperature based on the level of solar exposure. We can then compare this to the actual temperature record, and see how well the model fits the data set.

```
#Use this lm model to predict the values
MEL_weather_2019 <- MEL_weather_2019 %>%
  mutate(predicted_temp=predict(temp_model,newdata=MEL_weather_2019))

#Calculate the prediction interval
prediction_interval <- predict(temp_model,
                              newdata=MEL_weather_2019,
                              interval = "prediction")
summary(prediction_interval)

#Bind this prediction interval data back to the main set
MEL_weather_2019 <- cbind(MEL_weather_2019,prediction_interval)
MEL_weather_2019
```

Model fit is easier to interpret graphically. Let's plot the data with the model overlaid.

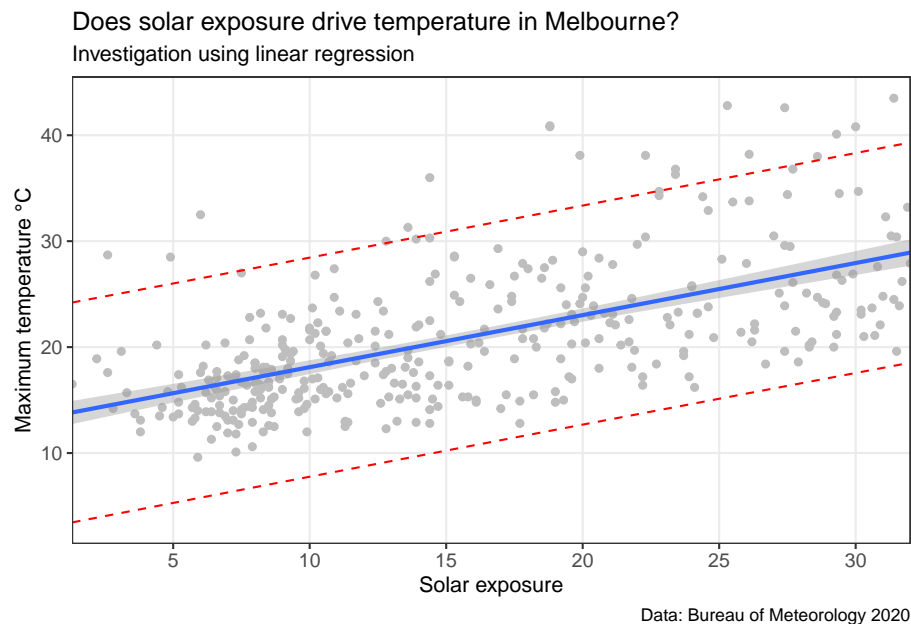
```
#Plot a chart with data and model on it
MEL_temp_predicted <-
ggplot(MEL_weather_2019)+
  geom_point(aes(y=Max_temp, x=Solar_exposure),
             col="grey")+
  geom_line(aes(y=predicted_temp,x=Solar_exposure),
            col="blue")+
  geom_smooth(aes(y=Max_temp, x= Solar_exposure),
              method=lm)+
  geom_line(aes(y=lwr,x=Solar_exposure),
            colour="red", linetype="dashed")+
  geom_line(aes(y=upr,x=Solar_exposure),
            colour="red", linetype="dashed")+
  labs(title =
        "Does solar exposure drive temperature in Melbourne?",
        subtitle = 'Investigation using linear regression',
        caption = "Data: Bureau of Meteorology 2020") +
  xlab("Solar exposure")+
  ylab("Maximum temperature °C")+

  scale_x_continuous(expand=c(0,0),
```

```
breaks=c(0,5,10,15,20,25,30,35,40))+

theme_bw()+
theme(axis.text=element_text(size=10))+
theme(panel.grid.minor = element_blank())

MEL_temp_predicted
```




This chart includes the model (blue line), confidence interval (grey band around the blue line), and a prediction interval (red dotted line). A prediction interval reflects the uncertainty around a single value (put simply: what is the reasonable upper and lower bound that this data point could be estimated at?). A confidence interval reflects the uncertainty around the mean prediction values (put simply: what is a reasonable upper and lower bound for the blue line at this x value?). Therefore, a prediction interval will be generally much wider than a confidence interval for the same value.

Analyse the residuals

```
#Add the residuals to the series
residuals_temp_predict <- MEL_weather_2019 %>%
  add_residuals(temp_model)
```

Plot these residuals in a chart.

```
residuals_temp_predict_chart <-  
  ggplot(data=residuals_temp_predict,  
    aes(x=Solar_exposure, y=resid), col="grey")+  
  
    geom_ref_line(h=0, colour="blue", size=1)+  
    geom_point(col="grey")+  
  
    xlab("Solar exposure")+  
    ylab("Maximum temperature (°C)") +  
    theme_bw() +  
    labs(title = "Residual values from the linear model")+  
  
    theme(axis.text=element_text(size=12))+  
  
    scale_x_continuous(expand=c(0,0))  
residuals_temp_predict_chart
```



Linear regression with more than one variable

The linear model above is *okay*, but can we make it better? Let's start by adding in some more variables into the linear regression.

Rainfall data might assist our model in predicting temperature. Let's add in that variable and analyse the results.

```
temp_model_2 <-  
lm(Max_temp ~ Solar_exposure + Rainfall, data=MEL_weather_2019)  
summary(temp_model_2)
```

We can see that adding in rainfall made the model better (R squared value has increased to 0.4338).

Next, we consider whether solar exposure and rainfall might be related to each other, as well as to temperature. For our third temperature model, we add an interaction variable between solar exposure and rainfall.

```
temp_model_3 <- lm(Max_temp ~ Solar_exposure +  
                    Rainfall +  
                    Solar_exposure:Rainfall,  
                    data=MEL_weather_2019)  
summary(temp_model_3)
```

We now see this variable is significant, and improves the model slightly (seen by an adjusted R squared of 0.4529).

Fitting a polynomial regression

When analysing the above data set, we see the issue is the sheer variance of temperatures associated with every other variable (it turns out weather forecasting is notoriously difficult).

However we can expect that temperature follows a non-linear pattern throughout the year (in Australia it is hot in January-March, cold in June-August, then starts to warm up again). A linear model (e.g. a straight line) will be a very bad model for temperature — we need to introduce polynomials.


For simplicity, we will introduce a new variable (Day_number) which is the day of the year (e.g. 1 January is #1, 31 December is #366).

```
MEL_weather_2019 <- MEL_weather_2019 %>%  
  mutate(Day_number=row_number())  
head(MEL_weather_2019)
```

Using the same dataset as above, let's plot temperature in Melbourne in 2019.


```
MEL_temp_chart <-
ggplot(MEL_weather_2019)+
  geom_line(aes(x = Day_number, y = Max_temp)) +
  labs(title = 'Melbourne temperature profile',
        subtitle = 'Daily maximum temperature recorded in Melbourne in 2019',
        caption = "Data: Bureau of Meteorology 2020") +
  xlab("Day of the year")+
  ylab("Temperature")+
  theme_bw()
```

MEL_temp_chart



We can see we'll need a non-linear model to fit this data.

Below we create a few different models. We start with a normal straight line model, then add an x^2 and x^3 model. We then use these models and the 'predict' function to see what temperatures they forecast based on the input data.

```
#Create a straight line estimate to fit the data
poly1 <- lm(Max_temp ~ poly(Day_number,1,row=TRUE),
             data=MEL_weather_2019)
summary(poly1)
#Create a polynomial of order 2 to fit this data
poly2 <- lm(Max_temp ~ poly(Day_number,2,row=TRUE),
             data=MEL_weather_2019)
```

```
summary(poly2)
#Create a polynomial of order 3 to fit this data
poly3 <- lm(Max_temp ~ poly(Day_number,3,raw=TRUE),
            data=MEL_weather_2019)
summary(poly3)

#Use these models to predict
MEL_weather_2019 <- MEL_weather_2019 %>%
  mutate(poly1values=predict(poly1,newdata=MEL_weather_2019))%>%
  mutate(poly2values=predict(poly2,newdata=MEL_weather_2019))%>%
  mutate(poly3values=predict(poly3,newdata=MEL_weather_2019))

head(MEL_weather_2019)
```

In the table above we can see the estimates for that data point from the various models.

To see how well the models did graphically, we can plot the original data series with the polynomial models overlaid.

```
#Plot a chart with all models on it
MEL_weather_model_chart <-
ggplot(MEL_weather_2019)+
  geom_line(aes(x=Day_number, y= Max_temp),col="grey")+
  geom_line(aes(x=Day_number, y= poly1values),col="red") +
  geom_line(aes(x=Day_number, y= poly2values),col="green")+
  geom_line(aes(x=Day_number, y= poly3values),col="blue")+


  #Add text annotations
  geom_text(x=10,y=18,label="data series",col="grey",hjust=0)+
  geom_text(x=10,y=16,label="linear",col="red",hjust=0)+
  geom_text(x=10,y=13,label=parse(text="x^2"),col="green",hjust=0)+
  geom_text(x=10,y=10,label=parse(text="x^3"),col="blue",hjust=0)+

  labs(title = "Estimating Melbourne's temperature",
        subtitle = 'Daily maximum temperature recorded in Melbourne in 2019',
        caption = "Data: Bureau of Meteorology 2020") +

  xlim(0,366)+
  ylim(10,45)+
  scale_x_continuous(breaks=
    c(15,45,75,105,135,165,195,225,255,285,315,345),
    labels=c("Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov")
    expand=c(0,0),
    limits=c(0,366)) +
```

```
scale_y_continuous(breaks=c(10,15,20,25,30,35,40,45)) +
xlab("")+
ylab("°C")+
theme_bw()+
theme(axis.text=element_text(size=12))+
theme(panel.grid.minor = element_blank())
```

MEL_weather_model_chart



Data: Bureau of Meteorology 2020

We can see in the chart above the polynomial models do much better at fitting the data. However, they are still highly variant.

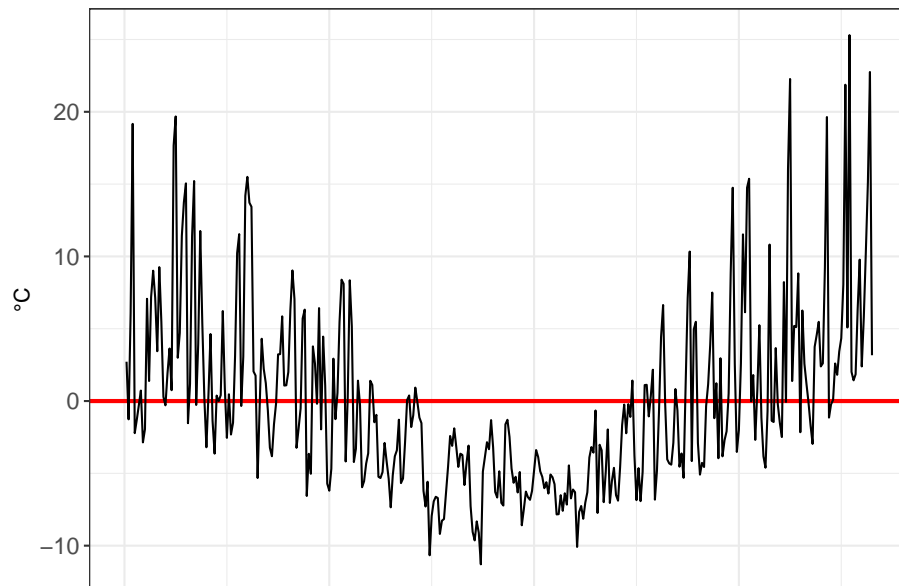
Just how variant are they? We can look at the residuals to find out. The residuals is the gap between the observed data point (i.e. the grey line) and our model.

```
#Get the residuals for poly1
residuals_poly1 <- MEL_weather_2019 %>%
  add_residuals(poly1)
residuals_poly1_chart <-
ggplot(data=residuals_poly1, aes(x=Day_number, y=resid))+
  geom_ref_line(h=0, colour="red", size=1)+
  geom_line()+
  xlab("")+
  ylab("°C")+
```

```

theme_bw()+
theme(axis.text=element_text(size=12))+
theme(axis.ticks.x=element_blank(),
axis.text.x=element_blank())
residuals_poly1_chart

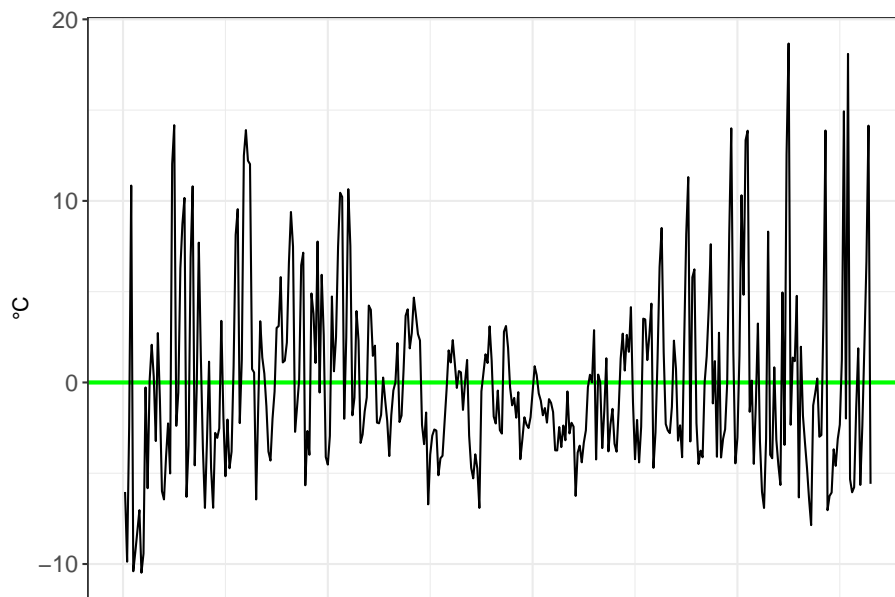
```



```


#Get the residuals for poly2
residuals_poly2 <- MEL_weather_2019%>%
  add_residuals(poly2)
residuals_poly2_chart <- ggplot(data=residuals_poly2,aes(x=Day_number, y=resid))+
  geom_ref_line(h=0,colour="green", size=1)+
  geom_line()+
  xlab("")+
  ylab("°C")+
  theme_bw()+
  theme(axis.text=element_text(size=12))+
  theme(axis.ticks.x=element_blank(),
axis.text.x=element_blank())
residuals_poly2_chart

```

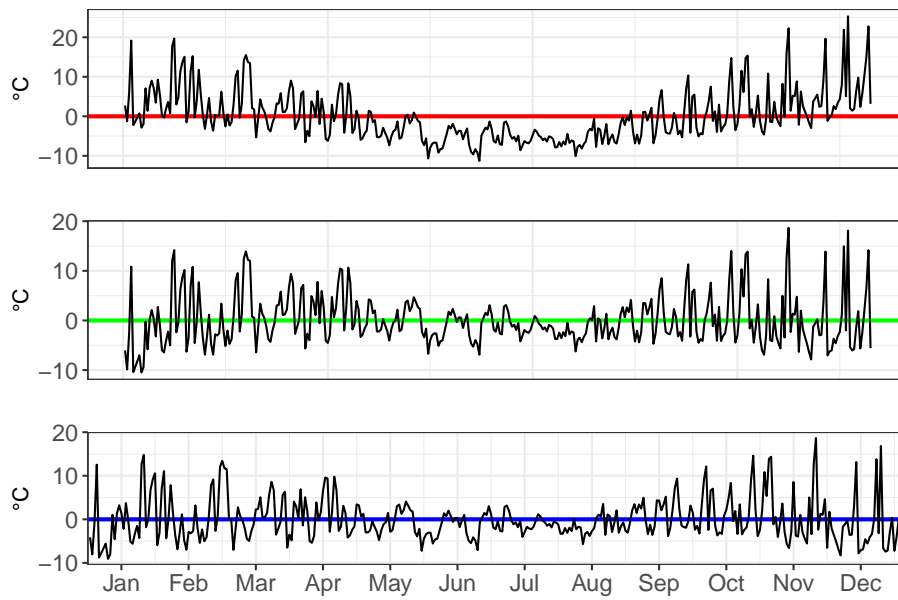


```
#Get the residuals for poly3
residuals_poly3 <- MEL_weather_2019 %>%
  add_residuals(poly3)
residuals_poly3_chart <- ggplot(data=residuals_poly3, aes(x=Day_number, y=resid))+
  geom_ref_line(h=0, colour="blue", size=1)+
  geom_line()+
  theme_bw()+
  theme(axis.text=element_text(size=12))+
  scale_x_continuous(breaks=
    c(15,45,75,105,135,165,195,225,255,285,315,345),
    labels=c("Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"),
    expand=c(0,0),
    limits=c(0,366))+
  xlab("")+
  ylab("°C")

residuals_poly3_chart
```



```
three_charts_single_page <- plot_grid(  
  residuals_poly1_chart,  
  residuals_poly2_chart,  
  residuals_poly3_chart,  
  ncol=1,nrow=3,label_size=16)  
three_charts_single_page
```



As we move from a linear, to a x^2 , to a x^3 model, we see the residuals decrease in volatility.

Hypothesis testing

Forecasting

Web scraping

Text mining

Geocoding

Drivetime analysis

Raster data

ABS economic indicators

Australian election data

Bayesian for the common man

